

Online Path Planning for Multi-Robot Coverage

Kevin Bradner

Department of Electrical Engineering and Computer Science

Case Western Reserve University

January, 2020

Contents

1	Introduction	1
1.1	Robotic Coverage Background	1
1.1.1	Spanning-Tree based Coverage Algorithms	3
1.1.2	Multi-Agent Robotic Coverage	8
1.2	Problem Statement	12
1.2.1	Rules for Robot Motion	14
1.2.2	Environment Description	15
1.2.3	Field of View and Information Sharing	18
1.2.4	Relation to Robotic Mapping	19
2	Simulation Development	23
2.1	Implementation of Problem Statement	24
2.1.1	Simulated Environment & Dynamics	24
2.1.2	Environment Interface	26
2.1.3	Environment Generation	29
2.2	Operation and Software Utilities	32
2.2.1	Running Simulations	32
2.2.2	Visualization Tools	33
2.3	Extra Features	33
2.3.1	Drone Dropout	33

3	Offline Planning Policies	35
4	Online Planning Policies	38
4.1	Persistent (Learning) Policies	38
5	Conclusion	40
5.1	Future Work	40
	Bibliography	42

List of Tables

List of Figures

1.1	Boustrophedon Decomposition	3
1.2	Spanning Tree Example	4
1.3	Spanning Tree Coverage Path	5
1.4	Neighbor Visiting Behavior for Online STC	6
1.5	Path Deformations in Scan-STC	8
1.6	MSTC Path Following Behavior	10
1.7	Multi-Robot Forest Coverage Path Creation	11
1.8	Odometry Errors and Corrected Map	20
2.1	Sample Environment Visualization	24

Acknowledgements

Abstract

Robotic coverage problems task one or more robots with the goal of visiting every location in a region. Algorithms that can perform this kind of task in a time efficient manner are useful for purposes such as mapping, cleaning, or inspection. This work considers a multi agent robotic coverage problem in which the shape of the region to be explored is known in advance, but additional information and challenges are discovered during task performance. A software package is created to simulate such a scenario, generate virtual environments to be covered, and interface with policy programs that command the robots. Offline path planning algorithms are developed, and their performance on this task is evaluated. Next, online variants of these algorithms are developed to respond to events and information encountered during task execution. It is shown that the online algorithms are more robust and better performing than their offline counterparts.

Chapter 1

Introduction

1.1 Robotic Coverage Background

Robotic motion planning problems are typically concerned with finding an achievable path between two robot configurations. These configurations could be the current location of a robotic vehicle along with its desired destination, the current pose of a robotic arm and the pose required for it to grasp an object, or any other pair of robot states that correspond to “start” and “goal” configurations. *Robotic coverage* problems are another kind of motion planning problem. The goal of a robotic coverage problem is to find a path for the robot to follow during which it will visit every location in a region. This region could be an area of floor to be vacuumed, the surface of a car body that needs to be painted, or the interior of a building that needs to be mapped. In all of these cases, the task is only complete once the entire region has been visited, or *covered* by a robot. Usually, another goal of these problems is to achieve coverage in a way that minimizes the time or number of movements required to complete the task. The robotic coverage problem can be viewed as a continuous generalization of the traveling salesman problem [2]. Thus, the general form of the problem is NP-Hard.

One particularly well studied type of robotic coverage problem is the case where the coverage region is some subset of the plane, and the robots are mobile robots that can only move in the plane. A good overview of this kind of coverage problem can be found in [3]. Algorithms to perform robot coverage in the plane have been studied since the late 1980's. Early algorithms for this purpose relied on the use of heuristics, and these algorithms usually could not guarantee that complete coverage would be achieved for all possible environments. By around 2000, several planar coverage algorithms were able to guarantee complete coverage for a variety of simple environment types. Many of these algorithms proved completeness by showing that the algorithm breaks down the coverage space into smaller regions that are relatively trivial to cover completely. Coverage algorithms that employ this kind of technique are known as *cellular decomposition* algorithms.

One type of cellular decomposition fits a square grid to the coverage space. Because a general arbitrary set in the plane can not be exactly expressed as a union of grid-aligned square regions, this technique is often called approximate cellular decomposition. Typically, each cell in this kind of decomposition can be covered entirely by a single robot configuration. For example, consider dividing a grass lawn into squares for a robotic lawn mowing task. If each square in the cellular decomposition is small enough, a lawnmower whose center is aligned with a square's center will have mowed that entire square of the lawn. In coverage problems where this assumption applies, proving that a coverage path visits each cell in the approximate decomposition is sufficient to guarantee complete coverage.

In addition to approximate cellular decomposition, there are also techniques that decompose a coverage space into a more diverse collection of cell shapes. Examples of such shapes include cells of fixed width whose top and bottom boundaries can vary in shape [11], or trapezoidal cells whose parallel sides all run in the same direction [4]. Depending on whether the union of these cells is exactly equal to

the coverage space, such techniques are called either *semi-approximate* or *exact* cell compositions. In the case of trapezoidal cell decomposition, it is possible to combine the trapezoidal cells into a smaller number of large cells, each of which are able to be covered by a simple back-and-forth sweeping motion. This approach is known as a *boustrophedon* decomposition [4].



(a) Two new cells created at critical point (b) Transition from two cells to one new one

Figure 1.1: Boustrophedon cell creation at critical points

1.1.1 Spanning-Tree based Coverage Algorithms

When coverage is performed by a square shaped tool attached to a mobile robot, and when the coverage area can be exactly decomposed into a grid of squares that have double the side length of the robot tool, it is possible to compute an optimal covering path in linear time using an approach based on spanning trees [8]. Gabriely and Rimon developed several variants of this approach with different time, memory, and prior knowledge requirements. In all cases, it was assumed that a mobile robot would complete the coverage task with a square shaped tool of size D . It was also assumed that the tool could only move in the four cardinal directions when completing the task. Finally, it was assumed that the region to be covered could be approximated by cells in a grid of squares with size $2D$. The authors argue that the final assumption is justified for realistic environments in which the size of the robot's tool is significantly smaller than the dimensions of the area to be covered. Under these assumptions, the *Spanning Tree Covering* (STC) algorithm is able to generate a path through the space which visits each grid-aligned square cell of size

D exactly once.




Figure 1.2: A small coverage environment and a spanning tree on that environment [8].

The offline variant of STC works by representing the coverage area as a graph G . Nodes on this graph represent the center of a square with size $2D$, and the set of nodes corresponds to the set of full cells when a grid with square size $2D$ is overlaid on the coverage space. Edges exist between any two nodes whose corresponding cells are adjacent in the coverage space. As the name suggests, STC's key step is the creation of a spanning tree on G , using any node in G as the root of the tree. Many algorithms exist to compute spanning trees, but two popular examples are based on breadth-first-search and depth-first-search [5].

Once this spanning tree has been created, the coverage area is subdivided into squares of size D . Each node of G coincides with the corners of four different squares in this subdivided map. Starting at any one of these squares, each of which is exactly the shape and size of the robot's coverage tool, the robot can simply move counterclockwise around the spanning tree as if it were performing wall following. In this way, each cell in the coverage space will be visited exactly once by the robot. Because of this property along with the fact that the path generated by STC ends in a position adjacent to where it started, the path found by STC can be considered a Hamiltonian cycle on G . Although constructing a Hamiltonian path on a general

planar grid is NP-complete [1], the assumed properties of G make it practical to compute such a path on this graph in linear time. The fact that this coverage path's start and end positions will be in adjacent cells can be beneficial in practice, as it allows deployment and retrieval of the robot to happen in approximately the same place.




Figure 1.3: Following the spanning tree of G achieves exact coverage and creates a cycle [8].

This algorithm has a few interesting variants. The first runs online, and effectively circumnavigates a depth first search spanning tree that it constructs on the fly [8]. Restriction to use of depth first search is the primary disadvantage of this approach compared to offline STC. While the coverage paths based on depth first search are equivalent to those created through any other method in terms of final path length, other spanning tree algorithms have practical benefits. For example, spanning tree algorithms such as Prim's algorithm or Kruskal's algorithm can create minimal spanning trees on weighted graphs. For details on Prim's algorithm and Kruskal's algorithm, refer to [5]. Setting the weights in a particular way can lead to coverage paths with desirable properties. For example, a minimal spanning tree on a graph where North-South edges are given less weight than East-West edges will tend

to move in long, parallel lines that minimize turning. In later work, Gabriely and Rimon develop an online approach called Scan-STC that achieves similar behavior. Scan-STC also makes additional improvements on the original STC as discussed below.

The second online variant of the STC algorithm is described by Gabriely and Rimon as *ant-like*. In this case, the robot must be able to mark visited sub-cells through some means such as leaving pebbles on the ground. In addition, the robot must be able to sense the presence of these makers in the nodes of G adjacent to its current position. By marking each visited sub cell as the robot leaves that cell, the robot is able to proceed on the same path it would have created in the depth first search based online approach. Because this algorithm uses markers to identify visited nodes, it requires only $O(1)$ memory. However, it requires the use of $O(N)$ markers, where N is the number of sub-cells in the coverage area. Nonetheless, this algorithm presents interesting ideas for multi-agent approaches and prevention of drift in pose estimation. In both of these cases, physical and immobile records of a robot's past location can inform the decision making of any robot currently near some of those records.



Figure 1.4: Online STC variants determine their behavior according to current position and whether neighbors have been visited before [8].

Finally, there is a variant of STC that achieves *exact* coverage of an environment, even when that environment can not be represented by completely filled cells on a grid of squares with size $2D$. However, this approach still requires that the connectivity of the environment is represented adequately by such a grid. The resulting algorithm simply notices any cells that are partially occupied by a previously unseen obstacle when it passes next to them in the course of normal online STC. When such a cell is encountered, a neighborhood is swept around the entire obstacle. Then, the robot returns to traversing the fully in-bounds cells according to its usual procedure. The ability to complete this sweeping motion is always guaranteed under assumptions used to model this problem. For details concerning the size of this neighborhood and the correctness of the resulting algorithm, see [8].

In the previously mentioned Scan-STC algorithm, Gabriely and Rimon improve STC to explicitly handle boundary cells, or graph cells that are partially occupied by an obstacle. Using the same notation as before, this algorithm generates a path with length no greater than $(n + m)D$, where n is the number of nodes in G , and m is the number of boundary cells [9]. The authors assume that no node will contain an arrangement of obstacles such that the graph of that node's sub-nodes is disconnected. If an obstacle exists such that the normal grid of nodes would violate this assumption, the corresponding location in the coverage area is assigned two different nodes. With that assumption in place, avoidance of obstacles is done by simply displacing the normal path that goes around the obstacle. This displacement tactic is what leads to the m term appearing in the maximum path length bound. The authors claim that in practical environments, the path length is closer to $n + \frac{m}{2}$.



Figure 1.5: Scan-STC adapts its motion in a variety of situations with partially occupied cells [9].

The other significant insight of Scan-STC relates to its name. By allowing the mobile robot to sense the state of the eight cells the surround its current position, it is possible to discover that one of the horizontal neighbors of the current cell has a free vertical neighbor, and that the current cell has a free vertical neighbor in the same direction. In this case, it is guaranteed that a path into the node's horizontal neighbor exists such that the vertical neighbors of these cells appear between the current cell and its horizontal neighbor. With this insight, it is possible to skip the creation of a horizontal edge between the current node and its horizontal neighbor. The benefit of this approach is that it leads to coverage paths that move primarily in the vertical directions. This can be useful for limiting the amount of turns performed, thereby reducing the number of opportunities to introduce dead reckoning error from imprecise turns.

1.1.2 Multi-Agent Robotic Coverage

Another imprint feature of a robotic coverage algorithm is whether that algorithm controls a single robot or a group. Multi-agent robotic coverage algorithms address

the latter case. Multi agent approaches to coverage can often complete a coverage task several times faster than a single agent approach because of the division of work that use of multiple robots enables. Use of multiple robots can also enhance robotic coverage by allowing for more precise localization through information sharing. Finally, multi-agent methods are often able to adapt to the failure of one or more robots, as there may be remaining robots capable of completing the coverage task [3]. As robotic hardware decreases in price, taking full advantage of these benefits is increasingly viable in terms of up-front costs. However, multi-agent techniques often require more complex motion planning strategies to coordinate the motion of multiple robots in a way that takes full advantage of these potential efficiency improvements and other benefits.

A large amount of research adapts the spanning tree approach to the case of multi-agent coverage [12][10]. In their work on this subject, Noam Hazon and Gal Kaminka generalize the STC algorithm to the multi agent case, resulting in an algorithm called *Multirobot Spanning-Tree Coverage*, or MSTC. The MSTC problem assumes that the coverage area can be adequately approximated by a grid of square cells with side length $2D$, as in the spanning tree coverage algorithm. As with the *original* version of STC, it is assumed that all cells in the coverage space are free of obstacles. It is also assumed that all robots have the same speed and tool size. A version of this assumption, that all robots are the same, applies to most of the work on multi-agent robotic coverage.

With these assumptions in place, the first version of MSTC simply constructs a spanning tree of the space offline, then has all of the drones start to follow the spanning tree from their initial positions, much the same as single agent STC. If the robots are able to share information about their status, i.e. working or broken, and whether they have completed their assigned section, then it is possible to have all working robots continue to follow the spanning tree path until the coverage task

is complete. Under these assumptions, complete coverage is guaranteed as long as one robot remains in working order until the end of the task.




Figure 1.6: The MSTC algorithm assigns portions of a shared loop to each drone based on its starting position [10].

If there are k robots exploring a coverage space with n sub-cells as described above, then the best case time for this algorithm is shown to be $\frac{n}{k} - 1$. This performance is achieved when the starting points of the robots are distributed such that each robot is evenly spaced along the length of the spanning tree coverage path. The worst case is stated to be $n - k - 1$, for the case where all k robots start on k adjacent sub-cells that appear together in the spanning tree graph. Implicit in both bounds is the assumption that no robots fail. In a pathological case with the same initial conditions used to give the previously stated worst case bound, the $k - 1$ robots that start directly in front of the last robot could all fail at time $n - k - 2$, just before the front robot reached the back robot's position. In this case, it would take another $n - 1$ steps for the remaining robot to complete the coverage task according to Algorithm 2 as described by the authors, leading to a total coverage time of $2n - k - 2$. The authors note that many practical multi-robot coverage scenarios require starting all robots from the same location, meaning that performance close to the worst

case bound $n - k - 1$ is relatively common. This motivates the development of a more sophisticated coverage algorithm. By assuming that backtracking is allowed, the authors develop an algorithm with a worst case of $\frac{n}{2} - 1$ for $k > 2$. However, because all algorithms in [10] only consider the possibility of moving forward and backward along the initially calculated spanning tree path, many opportunities for further improvement remain.

Work by Zheng et al. surpasses the average-case performance of MSTC with their *Multi-Robot Forest Coverage* (MFC) algorithm [12]. As the name suggests, MFC computes a separate spanning tree for each of the robots working on the coverage task. For convenience, the authors of MFC also assume that multiple robots can occupy the same sub-cell simultaneously, and that the coverage area can be represented by grid-aligned cells with side length twice that of the robot's coverage tool. Experimental results showed that MFC achieved performance close to a theoretical lower bound in many cases.



Figure 1.7: MFC assigns a different spanning tree to each robot.

In addition, Zheng et al. were able to prove that their algorithm's cover times are no worse than eight times the optimal value, thanks to properties of a tree cover algorithm for general graphs that they adapted to this purpose. The tree cover algorithm, developed by Even et al, finds a forest of $|R|$ trees that start from a set R of root positions and cover an entire graph $G = (V, E)$. This algorithm attempts to minimize the largest weight of its trees. The algorithm runs in polynomial time

and finds forests such that the maximum tree weight is no more than four times larger than optimal. Zheng et al. use a lightly altered version of this algorithm to determine the individual trees assigned to each robot in the space (Figure 1.7).

1.2 Problem Statement

As with most of the previously discussed research, this work presents multi-agent coverage algorithms based on spanning trees. The details of the coverage problem stated here vary somewhat from other work - these differences are motivated by a use case in which a team of aerial drones (UAVs, robots) must use cameras or other 2D imaging sensors to perform coverage of an area affected by disaster. Practical applications of this work would employ a combination of human effort and image recognition / processing techniques to identify people in need of aid, key structural damage that requires immediate repair, or other points of interest. Once any such point of interest has been identified and located, the relevant information can be used to guide response efforts.

Two particularly novel features are introduced to fit this setting:

1. Robots that are able to change the size of their field of view
2. Environments with subsections that require different levels of scrutiny

These two changes are closely related to one another. Specifically, the varying scrutiny requirements motivates the use of techniques that trade off between an observation's field of view and its level of detail. In turn, concrete motivation for an environment that requires varying levels of scrutiny is based on the idea that not all areas are worth viewing in great detail when gathering information for a search and rescue effort. For example, some areas of the coverage environment may consist of a fairly uniform surface such as a body of water or a large area of

flat, rocky terrain. Such areas may not be worth covering in detail because their uniform nature makes it easy to conclude that nothing of interest exists in these regions. At the opposite extreme, trying to observe a region with dense tree cover may not be practical in a coverage scenario constrained by an urgent timeline. In contrast, it may be worth looking closely at areas with partial occlusions, areas with visually complex damaged structures, or places where people are particularly likely to be in need of help.

As with the identification of actual people and other points of interest, the visual identification of which regions are likely to contain useful information that can be practically observed is abstracted away by this work. In its place, this work assumes access to a perfectly accurate sensor that describes the level of scrutiny required to view each environment cell in a drone's field of view. Another assumption of this problem is that the distribution of locations requiring one level of scrutiny or another is not known in advance of coverage. In practice, it may not be possible to know the exact location of collapsed structures, flooded regions, or other environment features that may appear without warning in an emergency.

The other significant novel feature of this problem formulation, robots with a variable field of view, arises naturally as a strategy to adapt to the varying scrutiny requirements of an environment. In this setting, the dimensions of a drone's field of view corresponds to the size of a robot's covering tool in a more traditional presentation of robotic coverage. This work assumes that a drone's field of view is always centered on a point directly below the drone, roughly corresponding to down-facing camera mounted on the bottom of a drone. Real drones can usually alter the size of a ground area being photographed. For example, increasing the size of the observed area may be achieved by ascending with a constant viewing angle or by using an optical camera zoom feature to increase the viewing angle. Likewise, decreasing the field of view may be achieved by descending, zooming in,

or by some combination of these.

Increasing the field of view has the obvious benefit of allowing a robot to cover more of its environment at once. Decreasing the field of view increases the detail with which a drone observes the area directly below it, and this allows high detail viewing of those areas for which a highly detailed view has been deemed useful. While the distribution of environment sections that require low vs high scrutiny is unknown a priori, it is assumed that the boundaries of the coverage area and of the obstacles that border it are known in advance.

1.2.1 Rules for Robot Motion

As previously stated, each drone in this formulation has access to moves that increase and decrease its field of view. For simplicity, it is assumed that a drone's field of view can be in only two states that correspond to two possible altitudes. A drone flying low can view a single cell in the environment's approximate cellular decomposition. A drone flying high can view a 3x3 grid of these environment cells. Depending on which altitude the drone currently occupies, it has access one of two vertical motions: ascend and descend. A somewhat arbitrary assumption of this work is that these vertical motions each take about the same length of time as moving cardinally between the centers of adjacent environment cells.

Unlike most other work on spanning-tree based coverage, this work assumes that the robots can move along diagonals, rather than only in the four directions orthogonal to the sides of its tool. For convenience of notation, the four directions commonly used by spanning tree robots will be referred to as *cardinal directions*, and the four directions that are offset by $\frac{1}{8}$ turn from the cardinal directions will be called *intercardinal directions*. In addition, concrete examples shown in this work may refer to a particular cardinal direction such as East, or a particular intercardinal direction such as South-West (SW).

One reasonable concern about the introduction of diagonal motion in a world represented by an approximate cellular decomposition is how it affects the connectivity of the coverage area. As with many coverage techniques that employ an approximate cellular decomposition of the environment, this work assumes that coverage paths will move directly from the center of one cell (node, patch, location) to the center of one of its neighbors. In addition, it is assumed that the robots are able to move directly between cells that share only a corner, even if one or both of the adjacent cells that would connect them are out of bounds. Since other work typically assumes that the robot is roughly the same size as the area it covers, this assumption requires justification. Because the problem explored here is motivated by a team of autonomous drones observing an area from above, it is safe to assume that the robots are significantly smaller than the area they are actively observing at any particular moment. In addition, the set of in-bounds cells is assumed to represent any cells that are completely covered by terrain to be observed, and so out of bounds cells may correspond to areas that are not *completely* obstructed. Finally, due to the ability of aerial drones to move in three dimensions, it is often possible for a drone to fly over any obstacles that aren't too close to the ground, even if these obstacles completely occlude the ground. As a result of these considerations, it seems justifiable to assume that the robots are able to move between diagonally connected in-bounds cells, as there is likely to be enough room for a drone to complete such a move in a realistic scenario.

1.2.2 Environment Description

In this work, the author assumes that all robots start out in the same location. It is possible that this assumption makes the work less general than other work that allows for arbitrary starting locations, such as the multi robot coverage work of Hazon and Kaminka [10]. However, starting all robots in one location is often a prac-

tical assumption, as it allows for easy deployment of robots from a single drop-off point. It is also worth noting that starting each robot in the same place often makes for a more challenging coverage task than starting with robots distributed throughout the space. This is because spatially distributed robots may already be close to evenly dispersed along a single Hamiltonian cycle through the coverage space, and so moving to this extremely convenient state may be a low cost operation. Note that in order to take advantage of this situation with the highest probability, robots must be able to move freely between their starting position and an idealized starting position for spanning tree following. This initialization step is not addressed by the MSTC algorithm or the MFC algorithm, and the MSTC algorithm in particular will often perform close to its worst case bound when it is necessary to deploy the robots near each other [12][10]. Most of the coverage algorithms developed in this work are also applicable to the case where each robot gets a different starting position. With these points in mind, assuming that all robots start in one location is unlikely to limit the usefulness of this work in practical applications.

Unlike many of the efforts to adapt spanning tree coverage to a multi agent case, this work allows for environments that can't be cleanly represented as a collection of obstacle-free cells twice the size of a robot's coverage footprint. This assumption complicates the issue of creating and traversing spanning trees significantly. However, it is likely to be an important practical assumption for many realistic tasks. For example, in a search and rescue setting, people may be trapped around collapsed structures and other difficult to navigate corners of the coverage area. The capability to plan spanning tree paths through such complex environments is also useful for the case where a robot has viewed an area with low scrutiny and must now revisit a small portion of that area that requires high scrutiny. Although it is not necessary to forbid revisiting locations that have already been adequately covered, a good path planning strategy should not visit these locations unless doing so is the most effi-

cient way of covering the locations that do require a visit. In addition, it may be known in advance that certain areas of the environment are simply not worth examining at all. For example, it is unnecessary for a drone to scout out areas where human search and rescue teams are already planning to visit due to ease of access or a likelihood of finding people there. As in the case of previously visited locations it is not strictly necessary to treat these areas as obstacles. However, any technique that can avoid visiting them without reducing the efficiency of its exploration of the rest of the environment will save time and perform better than if those areas were considered in-bounds. Finally, the drones that motivate this research have a field of view that is much larger than the tool sizes of robots used for cleaning, de-mining, or other popular applications of coverage. As a result, assuming that the environment contains only open spaces with at least four times this area is not realistic.

With these caveats in mind, the environment model used here is very similar to the models used in other research on spanning tree based coverage. It consists of a set of grid-aligned square cells that require coverage. This region is always finite, connected, and enclosed by a boundary that forms a single loop around the environment. As discussed previously, an arbitrary subset of the cells in an environment may be considered out of bounds. Because aerial drones are likely to occupy only a small footprint inside the three dimensional airspace over the relatively large region they are observing, it is assumed that multiple drones can occupy the same environment cell simultaneously without crashing. As discussed shortly, this does not present an issue from the perspective of being centered over the observed cell(s) because of information sharing.

1.2.3 Field of View and Information Sharing

Like much of the work that addresses multi-agent robot coverage, this work assumes that the robots share information while performing coverage. This assumption avoids a significant and complex barrier to implementing the described algorithms on real robots. However, because this work is motivated by a use case in which the robots are substantially smaller than their coverage area, and because an approximate cellular decomposition is used to describe the set of locations in the coverage area, it is possible to communicate full information between the robots in this scenario using only small and relatively infrequent messages. In addition, realistic flying robots are almost always equipped with long range communication devices to allow manual control by a human operator and to transmit observations in real time. For a search and rescue scenario in particular, it seems likely that the drones would be equipped with the ability to transmit high resolution photos at the same rate that coverage information is passed, in which case passing key planning information would require only a small fraction of the drone's available bandwidth. With these considerations in mind, it seems safe to assume that the level of communication required to share planning information between drones is at least possible in a realistic scenario.

Specifically, it is assumed that all drones have access to the following information at all times:

1. A map of the in bounds locations in the coverage region
2. The current location of every drone participating in the coverage task
3. A partially filled map in which any previously observed locations are marked and classified as requiring high or low scrutiny
4. Information about which previously seen locations were not observed in adequate detail and must be revisited from a low altitude

In addition, it is assumed that a central controller implements all motion planning, so a global multi robot motion plan is implicitly maintained at all times. For all intents and purposes, this information can also be thought of as shared.

Because this problem statement allows for more types of robot state changes than the previously described work, it is worth addressing the optimization criteria and cost function for this task. Much of the existing work on multi agent coverage aims to find a set of robot paths that complete the coverage task while minimizing the length of the longest path taken by any of the robots. In this work, it makes more sense to assign time costs to all of the moves available to the drones. Then, the optimization goal for this task is to achieve complete coverage in the least possible time. It is assumed that ascending and descending each take as long as a cardinal motion between adjacent environment cells. Intercardinal motion between cells that share a vertex is assumed to take a factor of $\sqrt{2}$ more time than cardinal motions in order to correspond with the greater euclidean distance between the centers of these cells.

1.2.4 Relation to Robotic Mapping

Whereas many applications of robotic coverage involve mechanical interaction with the environment being covered (e.g. vacuum cleaning, humanitarian mine removal), this work addresses a use case where the robots are observing an environment with cameras or other sensors. As a result, it bears some resemblance to the robotic mapping problem. Despite this meaningful similarity, this work focuses on simply visiting the entirety of an environment. The imagined application of this behavior in a search and rescue scenario is to capture anyone in need of help on camera. From there, it would be necessary to employ human effort or a complex image recognition tool to identify these people and send help. Note that providing this kind of information does not imply any of the other steps involved in creating an

accurate map of some area.

This focus on online path planning for abstract environments turns out to have very little to do with the main considerations of robotic mapping research. In a survey of the topic, Thrun provides this definition: "Robotic mapping addresses the problem of acquiring spatial models of physical environments through mobile robots" [13]. In a world of perfect robots with completely reliable sensors and actuators, solving this problem might be as simple as implementing the behavior of the online coverage algorithms discussed so far. Real robots are not perfect, though, and this is particularly true in applications where the robots must be inexpensive in order for the target application to be cost effective. As a result, approaching the robot mapping problem without accounting for noise in sensors and actuators leads to extremely bad performance. For example, odometry is one of the most widely used techniques for keeping track of a robot's position in space as it takes local measurements during a mapping operation. In principle, it is possible to determine a robot's precise location relative to some starting point by tracking all of its actuated motions and adding up the corresponding displacements in space. With noisy odometry, small errors in the robot's estimate of position and orientation can cascade, with one error affecting the accuracy of all of the measurements that follow. As the robot travels farther from its initial position, small errors in estimated orientation can be particularly damaging to the accuracy of the resulting map.



(a) Noisy odometry affects map quality



(b) Corrected map

Figure 1.8: Use of belief filtering and feature correspondence with loop closure [13].

In order to fix this odometry problem, it is often possible for mobile robots to use trackable shapes in their environment as navigation beacons. Assuming that the walls, doorways, and other measurable objects in the environment are stationary, these can be used as reference points by which to navigate and correct odometry errors. Whereas humans naturally use vision to guide their motion through the world, robots must be explicitly programmed to reconcile the incompatible information coming from odometry, sensors, and any other onboard information sources. Some of the most successful approaches to this reconciliation have relied on probabilistic techniques. By explicitly modeling the noise from various sources of information, it is possible to maintain a probabilistic *belief* about the state of the world and the robot's pose.

With clever implementations, these techniques can be both robust and relatively inexpensive from a computational standpoint. For example, Kalman filtering with Gaussian assumptions about sensor noise creates a loop in which a possible belief about robot state is repeatedly adjusted during actuation and then corrected based on sensor measurements. Due to some convenient mathematical properties of Gaussian distributions, this technique can combine signals to maintain an up to date belief with fairly efficient matrix and arithmetic operations. Even in applications where sensor noise is not exactly Gaussian, techniques like this often work well. Kalman filtering and related methods also rely on sensors and algorithms to perceive the robot's local environment, summarize this perception as a set of key features and their coordinates, and then repeatedly determine which of the features measured at one time correspond to a feature measured afterwards.

While all of these considerations are important for practical robotic mapping operations, they have little to do with the goal of discovering a time efficient multi agent path through an environment with the goal of achieving coverage. In addition, many realistic search and rescue scenarios involving drones may have reasonably

accurate compass and GPS signals onboard the drones. While these signals are by no means free of noise, they do not suffer from the cumulative error development that tends to occur when using odometry. In addition, a large amount of robot mapping research applies to mobile robots measuring feature rich indoor settings with sophisticated sensors such as lidar. Because drones are unlikely to carry a lidar or any comparably useful distance measurement sensor, many of the techniques and results from mapping don't apply here. Nonetheless, mapping research is likely to have considerable applicability in a realistic search and rescue scenario. This is particularly true if the environment being explored is largely unknown or significantly changed by a recent disaster. Thus, while the goals and methods of robotic mapping do not apply to this work, it would surely be necessary to combine mapping approaches with the techniques here in order to produce a holistically useful robotic answer to the problems of search and rescue.

Chapter 2

Simulation Development

This chapter describes a software package, *oprc_env*, that simulates the coverage scenario described in the introduction. This software exposes a standard interface, the *Policy* typeclass, that allows for interaction between the simulation and a wide variety of control algorithms. The package also includes functionality to generate large environment datasets based on custom mixture distributions of environment generation algorithms. *oprc_env* also includes functionality to run large batches of simulations with policies that are capable of modifying themselves between scenarios, thereby enabling the development and testing of online machine learning approaches to this problem. Finally, the package features a replay animation utility to aid in communicating results visually. In addition to describing these utilities, this chapter will specify many of the implementation details required to fully understand the coverage scenario and problem statement used in this research.

Note: names of datatypes, type constructors, and type constants in *oprc_env* start with capital letters, e.g. *Position*. For brevity, the names of these datatypes may be used throughout this chapter and the next in reference to both the computational abstractions and the concepts they model.



Figure 2.1: A sample environment (left) and the partial information available to robots exploring this environment (right).

2.1 Implementation of Problem Statement

2.1.1 Simulated Environment & Dynamics

The `oprc_env` package represents a multi agent coverage problem as a simulation in which all times and locations are represented with integer values. All ground positions in the environment are indexed by a pair of integers (x,y) . For ease of communication, the x and y coordinates are assumed to correspond to longitude and latitude, respectively. For example, the position $(2,1)$ is directly east of the position $(1,1)$, and the position $(0,0)$ is to the southwest of $(1,1)$. The minimum x and y coordinates in an environment are both 0 by default, although this is not

guaranteed to be true in general. All simulations start at time $T = 0$.

This design choice allows for simulations to be computed and performance to be evaluated more efficiently than if the environment was modeled as a continuous area or if the dynamics were modeled with continuous time. Since this work claims to achieve coverage paths that are applicable to continuous environments, as was the case with STC and other closely related algorithms, the use of a discrete simulation requires some justification. The work by Gabriely and Rimon on STC provides the inspiration and precedent justification for this choice [8]. Because all robots following STC move directly between the centers of adjacent cells, it is possible to encode the continuous path followed by an STC-controlled robot as a sequence of environment cells (or subcells, in keeping with the original work's terminology) visited. Since these cells are aligned with a square grid, it is trivial to show that these cells can be indexed by pairs of integer coordinates.

From there, assuming that time is divided such that all drone motions take an integer number of time steps allows for discrete time in simulation. While the ratio of distances corresponding to a cardinal and an intercardinal motion is technically irrational, it is perfectly reasonable to approximate this ratio as 10:14. As long as two drones in the same scenario don't use wildly different proportions of cardinal and intercardinal movements over any extended period, paths that take an equal amount of time in simulation will take a roughly equal amount of time in a continuous environment. Furthermore, the best performing algorithms to solve these scenarios use online techniques that would readily adapt to the occasional one time step delay of move completion. More details on these algorithms can be found in the remaining chapters.

The physical environment of each coverage scenario represents the locations that are in bounds, the level of detailed scrutiny required to cover each of those locations, and the positions of obstacles inside the environment. This is achieved

by a map from Position to Patch, where a Patch can represent an environment cell that requires observation at a particular level of detail. This level of detail can be either Close or Far. Patches that require Close detail must be observed by a drone flying at a Low altitude in order to be considered covered, while patches that require only Far detail will be covered once observed by a drone flying at either a High or Low altitude.

Each Patch in an environment can have up to eight neighbors. Four of these neighbors are adjacent to the patch in the four cardinal directions North, East, South, West, while the remaining four share a vertex with the patch and lie in one of the four intercardinal directions NE, SE, NW, SW. Given an in-bounds location, any arbitrary subset of these eight neighboring locations may be in bounds.

2.1.2 Environment Interface

At any given time during an in-progress coverage scenario, only a limited subset of environment information is made available to the algorithm that controls drone motions. The set of in bounds locations in the environment, also known as that environment's *Footprint*, is always known. The partial information available to an agent is represented by the EnvironmentInfo datatype, and this type contains a map from Position to PatchInfo. The set of Position values in this map's domain is normally the same Footprint that belongs to the real environment. A PatchInfo represents a state of knowledge about the patch at a particular location. Patches can be Unseen, Classified, or FullyObserved. As the names suggest, Unseen and FullyObserved correspond to patches that have never been seen at all and to patches that have been completely covered, respectively. A Classified patch is one where the level of observation detail required (DetailReq) is known, but the patch has not been fully covered yet. In practice, this only comes up for patches that require Close scrutiny and have only been seen by one or more drones flying at a High altitude,

in which case the patch would have only been observed with Far scrutiny.

The `EnvironmentInfo` datatype, together with the previously described `EnsembleStatus`, forms the content of the `WorldView` datatype. A `WorldView` represents all of the information about the current state of the simulation available to a control algorithm. Control algorithms plug in to the simulation via an instance of the `Policy` typeclass, which abstracts the behavior of any type capable of commanding directions to an `Ensemble` and potentially maintaining a notion of internal state. Specifically, the typeclass is defined as:

```
class Policy p where
  nextMove :: p -> WorldView -> (NextActions, p)
```

As the name suggests, `NextActions` is a datatype that expresses newly issued commands for some drones to follow. This is implemented as an association list:

```
type NextActions = [(Drone, Action)]
```

The `nextMove` function provided by any type `p` with an instance of `Policy` is called by the core simulation code to determine what moves have been commanded to each drone in the simulation at every time step. As suggested by the type signature, most sensible implementations of `nextMove` use a particular value of the policy type `p` to decide what moves to command next. A simple type for which a `Policy` instance would be easy to provide might contain a direct mapping from `WorldViews` to `NextActions`. In this case, an instance of `nextMove` would determine what value of `NextActions` to return by simply performing a lookup on this map at each time step. Along with this value, `nextMove` would return the original policy value along with the result of that map query.

Of course, an instance of `Policy` like the one described above is not easy to implement for anything but the smallest environments and number of drones. This is

because the number of possible WorldViews in an environment must be larger than the number of ways the drones can be distributed in that environment. It is easy to show that the latter notion corresponds to sampling from a bag of unique outcomes with replacement. Thus, for an environment with n cells and m drones, the number of WorldViews is much greater than n^m . This is already an inconvenient number of entries to store in a data structure, and each of these outcomes corresponds to a distribution of drones in space that is achievable. As a result, some number of entries for each of these would need to be included in any Policy implemented as a lookup table.

Counting the number of achievable WorldViews exactly, complete with all of the possible partial information states, is much more difficult. However, the established lower bound makes it clear that this quantity is not practical to handle in any direct way. This discussion may seem pointless for an implementation of Policy that is obviously unwise. However, getting a sense of this scale has implications that will inform the design of internal representations used by more effective Policy instances as discussed in later chapters. Luckily, because nextMove yields another value of p along with its NextActions decision, it is possible and often desirable for p to encode a control strategy that learns and plans, evolving with the context in which it currently operates.

Some control algorithms are capable of learning from experience and modifying behavior in a given simulation based on results from a previous one. The PersistentPolicy typeclass exists in order to provide the interface required for this functionality:

```
class Policy p => PersistentPolicy p where
  cleanup :: p -> (WorldView -> p)
```

The exact purpose of cleanup's type may not be immediately apparent. In short, the most useful policy representations tend to need access to some information

about the current environment before they can be initialized. The function type of `cleanup`'s output shows that it will return the right kind of value to make this initialization fit nicely with tooling developed for more general Policy instances. The fact that `cleanup` must be applied to a value of type `p` before it returns this value allows control algorithms to pass any encoded strategies learned during task execution to be applied to the completion of the next task. As later chapters will discuss, this capability is actually necessary for optimal performance when the distribution of environments is not known before task execution begins.

2.1.3 Environment Generation

In order to evaluate the coverage performance of any agent with a Policy instance, it is necessary to have a simulated environment for that algorithm to cover. It is also useful to have a large dataset of simulated environments for the agent to cover when collecting performance statistics. Finally, since some Policy instances may be capable of learning and specializing their behavior to the distribution of environments previously explored, it is useful to have fine control over the nature of the environments in a particular dataset. To meet all of these needs, *oprc-env* has a utility executable called *generate_environments* to create custom environments and environment datasets based on a specification of the distribution from which those environments should be drawn.

Before describing the capabilities of *generate_environments* in full, it is useful to consider the algorithms used to generate one specific environment. There are a few different algorithms for this purpose, and each one generates environments from a distribution that is qualitatively different from the others.

The simplest environment generation algorithm to describe is called the *BernoulliGen*. As the name suggests, this generator samples a required level of scrutiny from the same Bernoulli distribution for each location in the generated environ-

ment’s Footprint. As a result, most of the content of the *BernoulliGen* algorithm comes from the procedure used to generate Footprints. The *randomFootprint* function supplies this functionality.

The *randomFootprint* algorithm generates a Footprint, or set of in-bounds locations. While a Footprint alone does not constitute a usable coverage environment, it can be augmented to form a complete environment description by algorithms that assign a *DetailReq* value to each of the positions in the set. The type signature and corresponding parameter nicknames used by *randomFootprint* are as follows:

```
randomFootprint :: StdGen --a source of randomness
                  -> Int --allowed variation between border locations
                  -> Int -> Int -> Int -> Int --coordinate limits for Positions
                  -> Maybe Footprint

randomFootprint gen
                  varLimit
                  xMin xMax yMin yMax =
undefined --bulk of implementation excluded for brevity
```

Roughly, this function randomly selects one point on each face of the “bounding box” of legal position coordinates defined by the x/y min and max arguments. Then, interpolation is used to connect turn these points into a discrete approximation of a quadrilateral. Each coordinate on this perimeter is randomly perturbed by an amount whose magnitude must not exceed *varLimit*. Finally, the set of all positions inside this perimeter is computed and returned. The output type of a fully applied *randomFootprint*, *Maybe Footprint*, indicates that this function may fail to produce a valid Footprint. This is because the arguments for *randomFootprint* are sometimes sourced from configuration files or user input, and information from these sources may correspond to unrealizable requirements.

Given a random Footprint, the remainder of *BernoulliGen* simply assigns a *DetailReq* value to each location in the Footprint independently. The specific value assigned to each patch is determined by sampling a floating point number from the interval $[0, 1]$ and then comparing this value to a threshold provided as input to the function. The value of this threshold corresponds to the probability that any individual patch will be assigned the value *Close*, meaning that it must be observed by a Low drone before it is covered.

Other algorithms available for use with *generate_environments* produce an Environment in which the spatial distribution of *DetailReq* values features regions that are locally correlated in some way. A variety of these algorithms exist, and all of them rely on *randomFootprint* for the environment's basic shape. From there, however, these algorithms rely on procedures for filling in *DetailReq* values which are more sophisticated than the one used by *BernoulliGen*. In order to produce a somewhat structured output that still features randomness, many of these algorithms benefit from an efficient way of assigning *DetailReq* values to patches in a random order. This allows each new assignment to be drawn based on the presence of any other locations that have already been filled in.

The Fisher-Yates shuffle is an algorithm to reorder, or shuffle, the elements of a finite sequence. This algorithm provides a random reordering of a list of positions to support the many environment generation algorithms that require this functionality. It is also used for implementation of random drone failure and other simulation features that have not yet been described. It is possible to run this algorithm in-place on a sequence of length n in $O(n)$ time. In addition, given a way to randomly select an element from a range of integers with uniform probability, the Fisher-Yates shuffle has the useful property of producing any possible reordering of the sequence with uniform probability. The original work by Richard Durstenfeld to develop this algorithm for use on a computer can be found in [6]. The version presented here is

a generalization of that algorithm to operate directly on sequences of arbitrary data.

The *Shuffle* function assumes access to a function g that can accept an integer argument i and return an integer sampled uniformly from 0 to i , inclusive. It operates on l , a finite sequence of values of any datatype. For simplicity, it also assumes access to functions that get the length of a finite sequence and swap the elements of that sequence at two indices. Finally, note that sequence indexing starts at 0.

```
function SHUFFLE( $g, l$ )
```

```
   $n \leftarrow \text{length}(l)$ 
```

```
  while  $n > 0$  do
```

```
     $s \leftarrow g(0, n)$ 
```

```
     $\text{swap}(l, s, n)$ 
```

```
     $n \leftarrow n - 1$ 
```

```
  end while
```

```
  return  $l$ 
```

```
end function
```

The actual source code for this algorithm can be found in the FisherYatesShuffle module in *oprc_env*.

2.2 Operation and Software Utilities

2.2.1 Running Simulations

The Scenario type and the functions that operate on it are responsible for expressing and updating the state of a simulation. The type itself consists of a WorldState and a Policy for tracking the instantaneous state of the simulation, along with time and history parameters to allow simulation replays and performance analysis:

```

data Scenario p =
  Scenario {
    getPolicy :: p
    , getWorldState :: WorldState
    , getTime :: Integer
    , getHist :: MoveHistory
  }

```

The function that steps simulations forward in time takes any $(Policy\ p \Rightarrow Scenario\ p)$ and applies `nextMove` to that `Scenario`'s `Policy` and a `WorldView` generated from the current `WorldState`. The resulting policy is assigned to the next `Scenario` value. The `NextActions` portion of `nextMove`'s output is where most of the simulation's portion of the computation occurs. First, for replay and analysis, any newly commanded actions are added to the `MoveHistory`. In addition, time is incremented.

A function, *updateState*, figures out the new instantaneous `WorldState` from

2.2.2 Visualization Tools

2.3 Extra Features

2.3.1 Drone Dropout

In order to fully test the adaptability of `Policy` implementations, *oprc_env* includes functionality to run simulations with drone dropout. When this feature is in use, up to one drone may be permanently disabled at each time step in the simulation. Whether or not a drone drops out during a given time step is based on comparison of a random number sampled from the uniform distribution over the interval $[0, 1]$ and a customizable threshold value. It is not possible for more than one drone to fail during a single moment in time according to this formulation. Also, in order

to make sure that all scenarios are possible to complete, at least one drone is guaranteed to be operational at all times. These two assumptions and the basic model of drone dropout they apply to are not meant to be especially realistic. However, this model of dropout is sufficient for splitting Policy algorithms into a few coarse categories according to adaptability: those that fail when dropout occurs, those that complete a dropout scenario in an inefficient way, and those that smoothly adapt to each dropout event they occur. Note that animations of scenarios with dropout do not show disabled drones in a unique way. Rather, disabled drones are simply frozen in place with whatever state they had just before being disabled. Although this makes it theoretically difficult to distinguish between a disabled drone and one which is active but repeatedly receiving Hover commands, this issue does not come up in practice.

Chapter 3

Offline Planning Policies

It may not be immediately obvious that a strategy for switching between altitudes is helpful. The problem formulation used here is such that policies that only ever command drones to fly at Low altitude can still be guaranteed to complete coverage for all legal environments. However, this policy is not optimal for all environments. To understand this, consider the following toy instance of the coverage problem:

.....

The above example makes it seem like the "High Sweep First" approach may be the best choice for a policy to complete this task. However, this is also not the case for all environments. Consider the following environment, for which the relative performance of the same two policies from last time is flipped:

Both of these examples used very small environments in order to limit the length of explanations required. On this small a scale, the size and shape of the environment's Footprint provides a strong hint about the appropriate altitude switching strategy to employ when exploring the environment. Because the environment's Footprint is available to an agent before any decision needs to be made, it may be possible to create a policy that optimizes its behavior around that factor somewhat. However, for large environments that mostly consist of wide open space,

the importance of this contribution on a case by case basis is likely to be minimal. Setting aside considerations about Footprint shape, the different ratios of High vs Low scrutiny patches used in these examples provide the first hint that the a priori unknown content of each environment has implications for the best possible strategy. As an extremely simple proof that this is the case, consider the following tiny environments:

In both cases, assume that there is a single drone that starts at the top left position at a Low altitude.

In the first case, the best possible solution is to immediately Ascend. This action, whose move cost is as low or lower than all other move costs available in this simulation, completes the coverage task. Thus, this is clearly the optimal policy to use for exploring this environment in particular. In the second case, all environment cells require High scrutiny, and so must each be visited from a Low altitude. This means that the drone must move to each of the following states before the coverage task is complete:

1. Low over Position (1, 1)
2. Low over Position (0, 0)

In this case, one optimal policy would be to move East and then South-West. This policy causes the scenario to complete in 24 time steps. More importantly, any Policy that starts out with an Ascend would then need to Descend before any of the necessary observations to complete coverage can be performed. These two altitude switching moves require a total of 20 time steps to complete, making it clearly impossible to finish the task in less than the 24 time steps that another policy demonstrates. With that, it has been demonstrated that a policy's chance at optimally exploring either of these environments can be ruined at the *very first* move performed. The set of moves which keep a possibility of optimality in the first envi-

ronment, Ascend, and the set of moves that do the same in the second environment, MoveCardinal East, MoveCardinal South, are disjoint. All of this is true despite the fact that policies operating in these two environments must make a decision based on precisely the same initial information:

This proves that it is impossible to create a policy that makes optimal decisions without assuming some distribution of possible environments. The exact degree to which environment distributions will affect policy performance on larger environments is not addressed analytically in this work. However, as later sections will demonstrate, experimental results show that this difference is quite dramatic.

As discussed in the previous chapter, `oprc_env` is capable of generating quite a large number of different environments. Even if the exact parameters used by a particular environment generator are known in advance, it is not easy to characterize the distribution of resulting environments in a way that is useful for fine-tuning the design of policies. In addition, it seems intuitively likely that the distribution of real world environments in which this coverage behavior would be useful is similarly complex. This chapter will discuss the development and experimental optimization of Policy instances that do a good job on a complex distribution of environments.

Chapter 4

Online Planning Policies

4.1 Persistent (Learning) Policies

In the last chapter, it was proved that the optimal policy with which to perform coverage depends on the distribution environments that policy will be asked to solve. In that chapter, this problem was approached by hand crafting policies to perform as well as possible on a complex distribution of environments through experimentation and tweaking of parameters. In this section, some policy implementations are developed that can recreate a version of this experimentation and tweaking across multiple scenarios. Importantly, these policies perform these tweaks *automatically*, using the PersistentPolicy typeclass and associated environment interface tools to refine an approach specifically tuned to a distribution of environments encountered online.

It's worth addressing why this online learning behavior could be superior to a hand crafted approach. The distribution of real world environments in which you might want to deploy drone based robotic coverage is not at all easy to characterize. Furthermore, practical applications of this kind of technology may see it deployed to various owners around the world. Any one of these teams of robots is likely to

be deployed in environments with somewhat different properties the environments seen by other teams. As a result, there may be an opportunity to optimize the behavior of each of these teams of robots to the particular circumstances in which it is deployed. Such optimizations would not be easy to achieve through any sort of manual effort, so a well crafted policy that is capable of changing itself between individual operations may be the best way to achieve this optimization. In addition, there is some theoretical value in seeing how coverage policies can be automatically optimized for an experimentally determined distribution of environments. Much of the remainder of this chapter is devoted to exploring that topic.

Chapter 5

Conclusion

5.1 Future Work

Although this work is fairly permissive of complex environment shapes, it does not consider coverage of environment that can't be exactly represented as a square grid. As noted previously, the search-and-rescue scenarios that motivate this work could benefit from close inspection of tight corners and areas near the perimeter of environment boundaries. Since these boundaries may be due to collapsed or naturally occuring structures, their borders are unlikely to be aligned with cardinal directions. An algorithm that could adapt its behavior near these regions could be of great practical importance. For an example of other work that has achieved this kind of extension in the past, consider the exact coverage variant of STC as described in [8]. It may also be useful to represent the borders of regions requiring low or high scrutiny as arbitrary polygons, although the opportunity for improved algorithmic performance or other practical benefits are less clear in this case.

Another potential improvement of this work would allow for persistent coverage. Targets may be mobile in certain search and rescue situations, and so it is necessary to re-visit locations periodically in order to make sure nobody has moved

into them.

Bibliography

- [1] C. Papadimitrou A. Itai and J. Szwarcfiter. “Hamilton Paths in Grid Graphs”. In: *SIAM Journal on Computing* (Nov. 1982).
- [2] Esther M. Arkin, Sándor P. Fekete, and Joseph S. B. Mitchell. “The Lawn-mower Problem”. In: *CCCG*. 1993.
- [3] Howie Choset. “Coverage for robotics – A survey of recent results”. In: *Annals of Mathematics and Artificial Intelligence* 31 (Jan. 2001), pp. 113–126.
- [4] Howie Choset et al. “Exact Cellular Decompositions in Terms of Critical Points of Morse Functions”. In: *Proceedings of IEEE International Conference on Robotics and Automation (ICRA '00)*. Vol. 3. Apr. 2000, pp. 2270–2277.
- [5] Thomas H. Cormen et al. *Introduction to Algorithms*. 2nd. McGraw-Hill Higher Education, 2001. ISBN: 0070131511.
- [6] R. Durstenfeld. “Algorithm 235: Random Permutation”. In: *Communications of the ACM* (July 1964).
- [7] I. Endo. “Cooperative Sweeping by Multiple Mobile Robots with Relocating Portable Obstacles”. In: *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems* (Nov. 1996).

- [8] Y. Gabriely and E. Rimon. “Spanning-tree based coverage of continuous areas by a mobile robot”. In: *Annals of Mathematics and Artificial Intelligence* 31 (Oct. 2001).
- [9] Yoav Gabriely and Elon Rimon. “Competitive on-line coverage of grid environments by a mobile robot”. In: *Computational Geometry* 24.3 (2003), pp. 197–224. ISSN: 0925-7721. DOI: [https://doi.org/10.1016/S0925-7721\(02\)00110-4](https://doi.org/10.1016/S0925-7721(02)00110-4). URL: <http://www.sciencedirect.com/science/article/pii/S0925772102001104>.
- [10] N. Hazon and G. Kaminka. “Redundancy, Efficiency and Robustness in Multi-Robot Coverage”. In: *Proceedings of the 2005 IEEE International Conference on Robotics and Automation* (Apr. 2005).
- [11] Susan Hert, Sanjay Tiwari, and Vladimir Lumelsky. “A terrain-covering algorithm for an AUV”. In: *Autonomous Robots* 3.2 (June 1996), pp. 91–119. ISSN: 1573-7527. DOI: 10.1007/BF00141150. URL: <https://doi.org/10.1007/BF00141150>.
- [12] D. Kempe. “Multi-Robot Forest Coverage”. In: *IEEE/RSJ International Conference on Intelligent Robots and Systems* (Aug. 2005).
- [13] Sebastian Thrun. “Exploring Artificial Intelligence in the New Millennium”. In: ed. by Gerhard Lakemeyer and Bernhard Nebel. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc., 2003. Chap. Robotic Mapping: A Survey, pp. 1–35. ISBN: 1-55860-811-7. URL: <http://dl.acm.org/citation.cfm?id=779343.779345>.